

# Data2060 Final Project

Gaussian Naive Bayes for classification  
Never Converge

Zi Tao  
Jiaqing Dai  
Yingzhi Ma  
2025.12.10

# Introduction

- classification tasks involving continuous-valued features.
  - assume conditional independence between features.
  - Unlike the Bernoulli Naive Bayes models that assume discrete features, Gaussian Naive Bayes models each feature using a Gaussian (normal) distribution for every class.
- 
- mathematical simplicity, fast training time, and competitive performance

# Representation - Math

Given a feature vector

$$\mathbf{x} = (x_1, x_2, \dots, x_d),$$

GNB models the conditional likelihood for each class  $y \in \{1, \dots, K\}$  as:

$$P(x_i | y) = \mathcal{N}(x_i; \mu_{y,i}, \sigma_{y,i}^2) = \frac{1}{\sqrt{2\pi\sigma_{y,i}^2}} \exp\left(-\frac{(x_i - \mu_{y,i})^2}{2\sigma_{y,i}^2}\right).$$

Under the independence assumption, the joint likelihood becomes:

$$P(\mathbf{x} | y) = \prod_{i=1}^d P(x_i | y).$$

Using Bayes' rule, the posterior is:

$$P(y | \mathbf{x}) \propto P(y) \prod_{i=1}^d P(x_i | y).$$

To avoid floating-point underflow, we compute the \textbf{log-posterior}:

$$\log P(y | \mathbf{x}) = \log P(y) + \sum_{i=1}^d \log \mathcal{N}(x_i; \mu_{y,i}, \sigma_{y,i}^2).$$

The predicted class is:

$$\hat{y} = \arg \max_y \log P(y | \mathbf{x}).$$

Thus, the feature representation converts continuous inputs into probabilities governed by class-specific Gaussian distributions.

Bayes Rule:

$$P(A | B) = \frac{P(B | A) \cdot P(A)}{P(B)}$$




Image source: Pierce, Rod, 2025, 'Normal Distribution, Math Is Fun,  
Available at:  
<<https://www.mathsisfun.com/data/standard-normal-distribution.html>>.

# Loss - Maximum Likelihood Estimates

Class priors:

$$P(y) = \frac{N_y}{N}$$

Feature mean:

$$\mu_{y,i} = \frac{1}{N_y} \sum_{x \in y} x_i$$

Feature variance:

$$\sigma_{y,i}^2 = \frac{1}{N_y} \sum_{x \in y} (x_i - \mu_{y,i})^2$$

At evaluation time, models are typically assessed using classification accuracy or cross-entropy loss:

$$\text{Loss} = - \sum_{n=1}^N \log P(y_n | x_n)$$

# Optimizer

GNB does not require iterative optimization. All parameters have closed-form MLE solutions, making the algorithm extremely efficient.

Parameters learned:

Class prior probabilities  $P(y)$  Feature means  $\mu_{y,i}$  Feature variances  $\sigma_{y,i}^2$

# Numerical Techniques

1. Using log-space to prevent numerical underflow

$$P(x | y) = \prod_{i=1}^d \mathcal{N}(x_i | \mu_{y,i}, \sigma_{y,i}^2) \quad \rightarrow \quad \log P(x | y) = \sum_{i=1}^d \log \mathcal{N}(x_i | \mu_{y,i}, \sigma_{y,i}^2)$$

2. Variance smoothing: adding  $\epsilon$  to the variance:

$$\sigma_{y,i}^2 = 0 \quad \rightarrow \quad \sigma_{y,i}^2 \leftarrow \sigma_{y,i}^2 + \epsilon$$

3. Closed-form MLE for comp

$$\mu_{y,i} = \frac{1}{N_y} \sum_{x \in y} x_i$$

$$\sigma_{y,i}^2 = \frac{1}{N_y} \sum_{x \in y} (x_i - \mu_{y,i})^2$$

# Numerical Techniques

Train:

initialize priors  $P(y)$ , means  $\mu[y][i]$ , variances  $\sigma^2[y][i]$   
for each class  $c$  :

$X_c \leftarrow$  all rows of  $X$  where label =  $c$

$N_c \leftarrow$  number of samples in class  $c$

$P(c) \leftarrow N_c/N$

for each feature  $i \in \{1, \dots, d\}$  :

$\mu[c][i] \leftarrow$  mean of  $X_c[:, i]$

$\sigma^2[c][i] \leftarrow$  variance of  $X_c[:, i] + \epsilon$  (smoothing)

return priors, means, variances

Predict:

To predict a class for a given input  $\mathbf{x}$ :

for each class  $c$  :

$\text{log\_post}[c] \leftarrow \log P(c)$

for each feature  $i \in \{1, \dots, d\}$  :

$\text{log\_post}[c] += \log \text{Gaussian}(x[i] | \mu[c][i], \sigma^2[c][i])$

return class with maximum  $\text{log\_post}$

# Compare With sklearn GaussianNB

# Motivation


- We implemented Gaussian Naive Bayes from scratch
- We need to verify that our math and implementation are correct
- sklearn's GaussianNB is a trusted reference model
- Matching sklearn confirms correctness and numerical stability

# Dataset Introduction

Public Dataset from Kaggle: [Naive-Bayes-Classification-Data.csv](#)

The dataset contains 995 entries and 3 columns:

- glucose: numeric feature (blood glucose level)
- bloodpressure: numeric feature (blood pressure)
- diabetes: binary label (0 = no diabetes, 1 = diabetes)



NaKrani, H. (2023). Naive Bayes Classification Data  
[Dataset]. Kaggle. <https://www.kaggle.com/datasets/himanshunakrani/naive-bayes-classification-data>

# Data Preparation & Experiment Setup

```
import pandas as pd
from sklearn.model_selection import train_test_split

df = pd.read_csv('Naive-Bayes-Classification-Data.csv')
df.head()

X = df[['glucose', 'bloodpressure']].values
y = df['diabetes'].values

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)
```

- Selected continuous features: glucose, blood pressure
- Target label: diabetes (0 or 1)
- Train-test split: 70% / 30%
- random\_state = 42 for reproducibility
- Used stratify=y to keep class balance

# Results: Our Model vs sklearn GaussianNB

```
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

# Train our model
model_my = GaussianNaiveBayes()
model_my.train(X_train, y_train)

# Predict with our model
y_pred_my = model_my.predict(X_test)
acc_my = accuracy_score(y_test, y_pred_my)

# Train sklearn model
model_sk = GaussianNB()
model_sk.fit(X_train, y_train)

# Predict with sklearn model
y_pred_sk = model_sk.predict(X_test)
acc_sk = accuracy_score(y_test, y_pred_sk)

# Compare results
print("My model accuracy:      ", acc_my)
print("Sklearn model accuracy: ", acc_sk)
print("Predictions identical:   ", np.array_equal(y_pred_my, y_pred_sk))
```

```
My model accuracy:      0.9264214046822743
Sklearn model accuracy: 0.9264214046822743
Predictions identical:   True
```

- Trained both models on the same training split
- Tested both on the same test split
- Our model accuracy: **0.9264**
- sklearn accuracy: **0.9264**
- Predictions were identical sample-by-sample

# Why Do the Results Match?

- Gaussian NB uses closed-form estimates for parameters
- Both implementations compute same:
  - class priors
  - means
  - variances
- Both use log-likelihood for stability
- So identical predictions are expected when math is correct

# Summary

- Interesting
  - Gaussian NB has very clean math
    - Parameters have closed-form formulas (no optimization)
    - Only needs simple statistics (means & variances)
    - Extremely fast to train
- Challenges
  - Equations

Thank You